

# blocking

## uniprocessor blocking

When a task of lower base-priority is executing instead of a higher priority one.

How blocking impacts schedulability?

- s-oblivious: blocking counts as execution. Safe but unnecessarily pessimistic. Most schedulability test can use this technique. Example:

$$\sum_{\tau_i} \frac{e_i + B_i}{p_i} \leq 1$$

- s-aware: blocking does not count as execution. Safe and tight. Few schedulability test can use this technique. Example:

$$B_j + \sum_{\tau_i} \frac{e_i}{p_i} \leq 1 \quad \forall \tau_j$$

# blocking in multiprocessor

uniprocessor blocking  $\neq$  multiprocessor blocking

- processors *can idle*
- access to resources is *parallel*

What to do while waiting for a locked resource?

- suspend: let other tasks execute
- spin: waiting and holding cpu

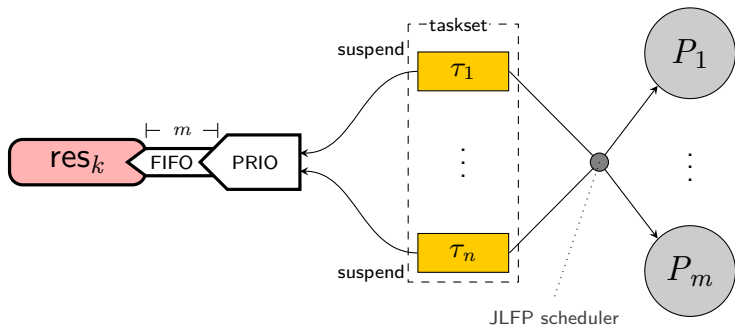
Not obvious which is best in multiprocessor

- changing priority or suspending blocked tasks *does not speed up* the release of resources (priorities *might not* have the same meaning across processors)
- spinning *does waste* cpu without letting any task progress

## $\mathcal{O}(m)$ Locking Protocol

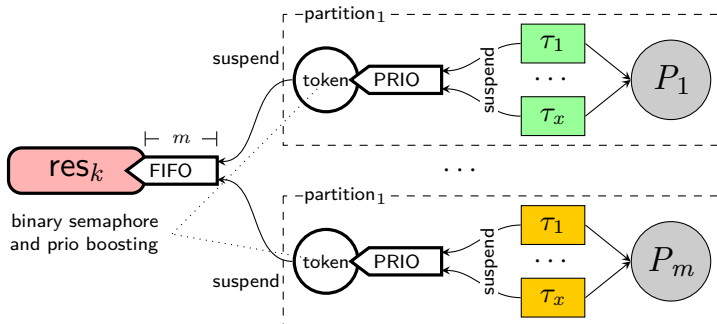
- suspension-based protocol for global and partitioned algorithms
- using queues to achieve optimal  $\mathcal{O}(m)$  blocking
  - ▶ FIFO queue: serializing access to shared resource, preventing starvation
  - ▶ PRIO queue: speeding up higher priority tasks (lower prio tasks are not blocked if higher prio tasks are suspended)
- blocking term usable inside s-oblivious schedulability test
- nested resources only with group locks

# OMLP - global



- blocking suffered only by tasks using resources
- per-request blocking is  $b_k = 2(m - 1)\omega_k$ ,  $\omega_k$  length of max critical section for  $res_k$
- all resources are global resources

# OMLP - partitioned



- limiting access to global resources: per-partition *contention token*. Must be acquired before requesting *any* global resource (token + PRIO queue shared for all global resources)
- releasing resources as soon as possible: *priority boosting* for tasks queued in global resources (at most 1 per partition)

# OMLP - partitioned

- three kinds of blocking:

- ①  $b^{prio}$ : caused by priority boosting (**any task**)

$$b^{prio} = \max_k \{\omega_k\}$$

- ②  $b_k^{fifo}$ : caused by waiting in FIFO queue (only if using  $res_k$ )

$$b_k^{fifo} = (m - 1)\omega_k$$

- ③  $b^{trans}$ : caused by the contention token (if using any global res)

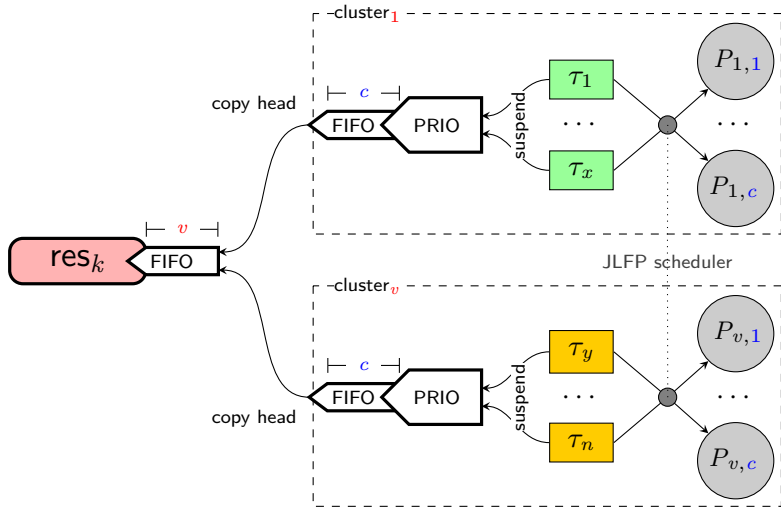
$$b^{trans} = (m - 1)\max_k \{\omega_k\}$$

- tasks suffer (extensively) from unrelated critical sections
- some resources can be local (using ICPP/SRP)

## $\mathcal{O}(m)$ Independence-preservation Protocol

- for clustered algorithms
- suspension-based protocol
- generalizes OMLP-glob and OMLP-part (as clustered for global and partitioned schedulers)
- avoids shortcomings from OMLP-part by requiring *intra-cluster migration*
  - ▶ theorem: intra-cluster migrations are necessary to not suffer from unrelated critical sections
- blocking term usable inside s-oblivious schedulability tests

# OMIP

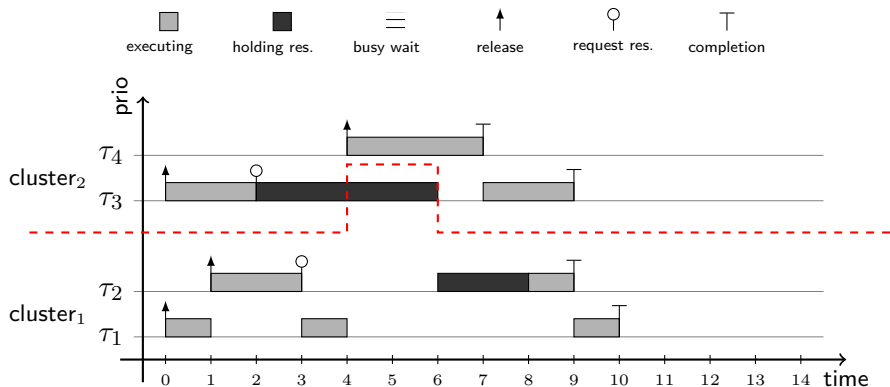




- head of per-cluster FIFO queue participates in global FIFO queue
- each global resource has a private per-cluster FIFO+PRIO queues
- head of global FIFO queue can migrate and inherit priority of other tasks queued in global FIFO queue
- per-request blocking  $b_k$  (only if using  $\text{res}_k$ )

$$b_k = (2m - 1)\omega_k$$

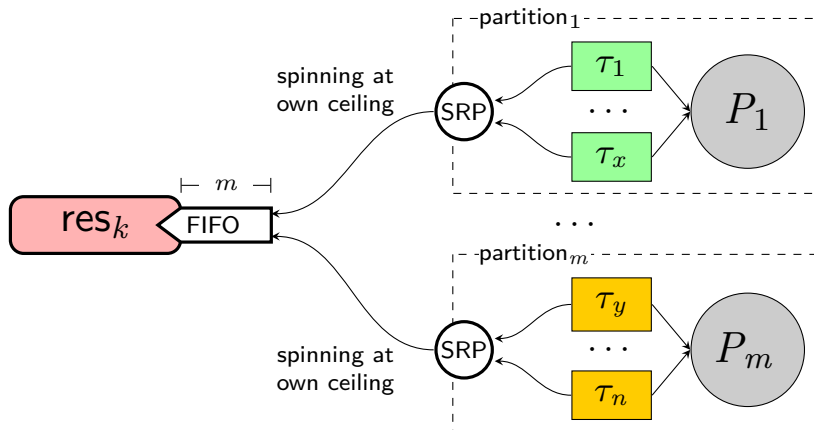
# OMIP - runtime example



- $t = 3$ : task  $\tau_2$  suspends and task  $\tau_1$  resumes execution
- $t = 4$ : task  $\tau_3$  migrates to cluster<sub>1</sub> and preempts task  $\tau_1$

## *Multiprocessor resource sharing Protocol*

- for partitioned algorithms
- spinning-based protocol
- generalizes uniprocessor RTA
- assumes availability of helping mechanism
  - ▶ *task migration approach*: task migrates and executes in place of the “helper”
  - ▶ *duplicated execution approach*: assuming resources have internal status and their use does not produce side effects
- nesting by:
  - ▶ using resources always in the same order (avoid circular wait) and ad-hoc analysis
  - ▶ group locks



- worst-case resource usage must consider parallelism

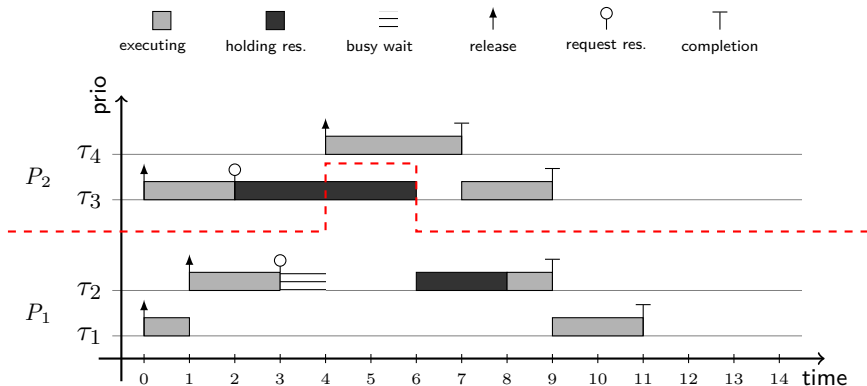
$$\hat{\omega}_k = m\omega_k \Rightarrow \hat{e}_i = e_i + \sum_{\text{res}_k \in \tau_i} (m-1)\omega_k$$

- access to global resources through local SRP
- waiting for locked global resources by *spinning at local ceiling* (remaining preemptable)
- head of global FIFO queue, if preempted, executes in place of other queued spinning tasks
- per-partition RTA equation

$$R_i = \hat{e}_i + B_i + \sum_{hp(\tau_i)} \left\lceil \frac{R_i}{p_j} \right\rceil \hat{e}_j,$$

$B_i$  = uniprocessor SRP-like blocking term using  $\hat{\omega}_k$

# MrsP - runtime example



- $t = 3$ : task  $\tau_2$  start spinning at ceiling priority
- $t = 4$ : task  $\tau_3$  migrates to  $P_1$  and executes in place of  $\tau_2$