



MrsP

Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments

# Multiprocessor resource sharing Protocol

## Implementation and evaluation

Sebastiano Catellani

University of Padua  
Supervisor: Prof. Tullio Vardanega



# Overview

MrsP

Sebastiano  
Catellani


MrsP

Proposed  
solution

Implementation

Experiments

- 1 MrsP
- 2 Proposed solution
- 3 Implementation
- 4 Experiments



**Figure:** Partitioned Fixed-Priority scheduler on a platform with  $m$  processors ( $P_1, \dots, P_m$ ) and a global resource



Burns and Wellings design a multiprocessor extension of PCP/SRP with the aim of adapt a schedulability analysis to the protocol

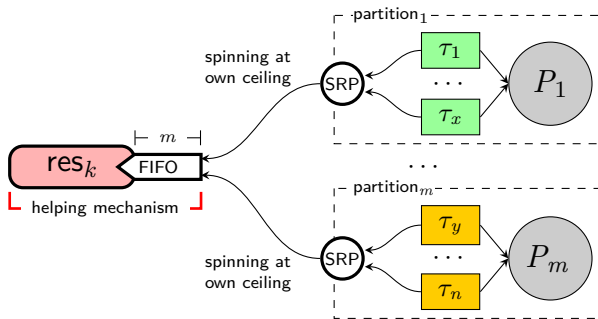
### Response Time Analysis incorporating PCP/SRP

The parameter  $e_j$  reflects the **contention** for the resource ( $r$ ):

$$e_j = |\text{map}(G(r))| \times c_j$$

$$R_i = C_i + \max\{e_j, \hat{b}\} + \sum_{\tau_j \in hp(i)} \lceil \frac{R_i}{T_j} \rceil C_j$$

$$C_i = WCET_i + \sum_{r^j \in F(\tau_i)} n_i e_j$$



### Protocol's properties

- It inherits the properties of PCP/SRP
- At most one job per processor requires the resource
- The length of the requests queue is at most  $|map(G(r_j))|$
- At most  $e_j$  to gain the resource and to execute the critical section



# Proposed solution

## Algorithm

- 1) Each resource has a set of ceilings, one for each processor
- 2) An access request causes the rise of the job's priority and activates a local ceiling
- 3) The requests are queued and served in arrival order
- 4) A job executes, until resource's release, at the inherited priority
- 5) If preempted, the lock holder migrates to the first processor available

### Key features

- Points 2 and 4 make MrsP **independence-preserving**
- Point 5 guarantees a **limited waiting and blocking time**



# Implementation

## Data structures

MrsP

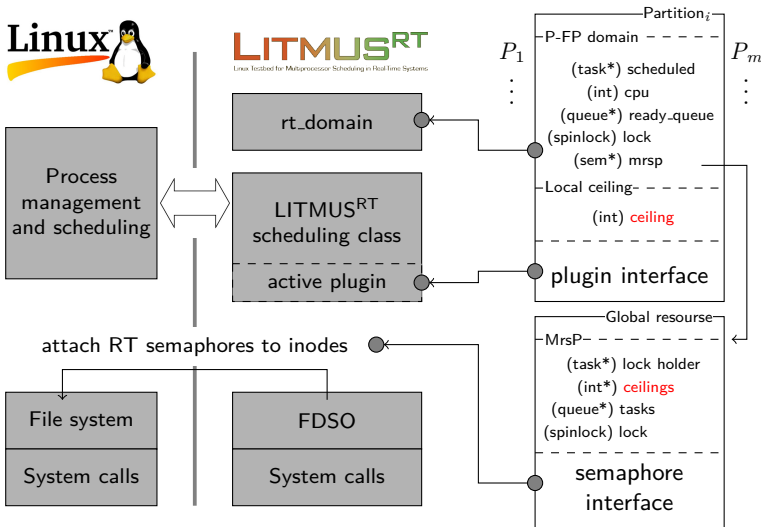
Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments





# Implementation

## Queue management - 1

MrsP

Sebastiano  
Catellani


MrsP

Proposed  
solution

Implementation

Experiments

Focused on managing the access requests queue



If preempted, the lock holder ( $J_x$ )

- 1 inherits the ceiling ( $c_3 + 1$ )
- 2 migrates to  $P_3$
- 3 preempts  $J_y$





# Implementation

## Queue management - 2

IMPP

Sebastiano  
Catellani


MrsP

Proposed  
solution


Implementation

Experiments


The job will be re-queued in the *ready\_queue*



The algorithm catches the operations that



(a) a processor becomes available



(b) add a new request to the queue



# Implementation

Primitive: `mrsp_lock`

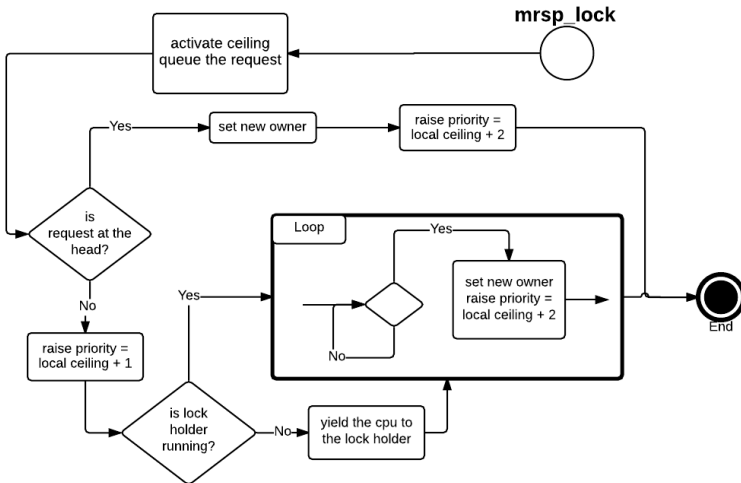
Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments





# Implementation

## Primitive: mrsp\_unlock

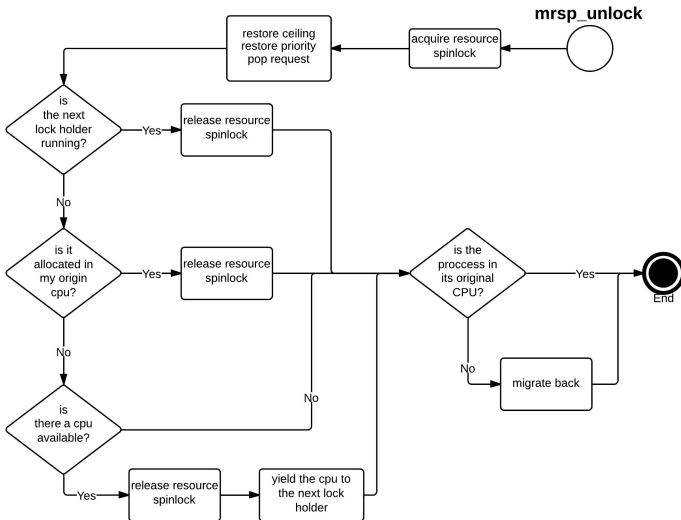
Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments





# Implementation

Primitive: pfp\_schedule and finish\_switch - 1


Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments



- $t_1$ :  $J_2$  is marked for migration
- $t_2$ :  $J_4$ 's priority is lower than the local ceiling
- $t_3$ : default migration mechanism



# Implementation

Primitive: pfp\_schedule and finish\_switch - 2


Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments



- $t$ :  $J_1$  completes and  $P_1$  returns available



# Experiments

## Overview

MrsP

Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments

### Experiment #1: Comparison among protocols

MrsP outperforms protocols based on simple ceiling or non preemption

### Experiment #2: Sampling of the overheads

MrsP brings benefits at reasonable costs

### Experiment #3: Absence of global resources

The protocol doesn't interfere with the scheduler



# Experiment #1

## Comparison among protocols - 1

MrsP

Sebastiano  
Catellani

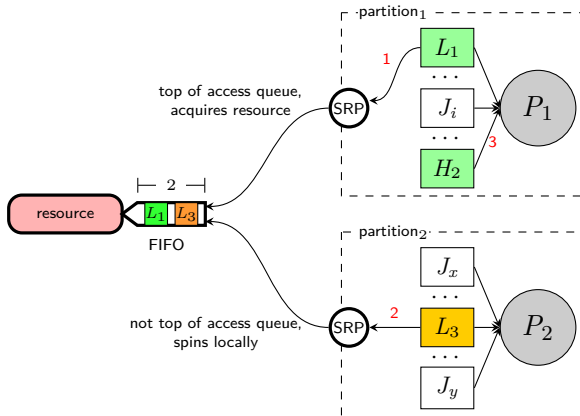
MrsP

Proposed  
solution

Implementation

Experiments

The experiment observes the **response times** of  $L_1$ ,  $H_2$  and  $L_3$  while varying the **critical section length** and the **WCET** of  $H_2$





# Experiment #1

## Comparison among protocols - 2

MrsP

Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments




Figure: Response time of  $L_1$





# Experiment #1

## Comparison among protocols - 3

MrsP

Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments

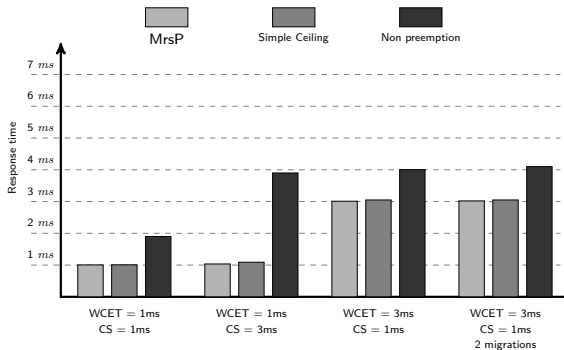


Figure: Response time of  $H_2$



# Experiment #1

## Comparison among protocols - 4

MrsP

Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments




Figure: Response time of  $L_3$



# Experiment #2

## Sampling of the overheads

MrsP


Sebastiano  
Catellani

MrsP


Proposed  
solution

Implementation


Experiments



(a) `mrsp_lock`



(b) `mrsp_unlock`



(c) `finish_switch`



# Experiment #3

## MrsP without global resources

IMRP

Sebastiano  
Catellani

MrsP

Proposed  
solution

Implementation

Experiments

The collected data show the same number of deadline miss




Figure: Number of *deadline miss*



# Experiment #3

MrsP without global resources - pfp\_schedule performance

MrsP


Sebastiano  
Catellani

MrsP


Proposed  
solution

Implementation


Experiments



(a) pfp\_schedule: Min



(b) pfp\_schedule: Max



(c) pfp\_schedule: Average