



MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion

Multiprocessor resource sharing Protocol

Implementation and evaluation

Sebastiano Catellani

University of Padua
Supervisor: Prof. Tullio Vardanega

9 October 2014



Overview

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion

- 1 Introduction
- 2 MrsP
- 3 Proposed solution
- 4 Implementation
- 5 Experiments
- 6 Conclusion



Real-Time System Model

Workload

Workload

- $T = \{\tau_1, \dots, \tau_n\}$
- $\tau_i = (C_i, P_i, D_i)$

A task τ_i is characterized by a utilization factor $U_i = C_i/P_i$




Figure: Task τ_i releases a job J_i in response to a triggering event



Real-Time System Model

Resources and scheduling algorithm

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion

Resources

- **Processors**, necessary to execute
 - ▶ Assigned by a scheduling algorithm
- **Memory, lock, I/O, ...**, provide functionality
 - ▶ Need mutual exclusion access → resource access protocol
 - ▶ Can be *local* or *global*

Scheduling algorithm

It determines which task has to be executed on each processor at any time




Real-Time System Model

Scheduler categorizations

MrsP

Sebastiano
Catellani

By the **number of ready tasks queues**:



and by the **priority assignment function**:

- *Fixed-Priority*
- *Job-Level Fixed-Priority*
- *Job-Level Dynamic-Priority*



Real-Time System Model Platform

MrsP

Sebastiano Catellani

Introduction

MrsP

Proposed solution

Implementation

Experiments

Conclusion




Figure: Partitioned Fixed-Priority scheduler on a platform with m processors (P_1, \dots, P_m) and a global resource



MrsP

Sharing resource in multiprocessor systems - 1

Sebastiano
Catellani

Introduction

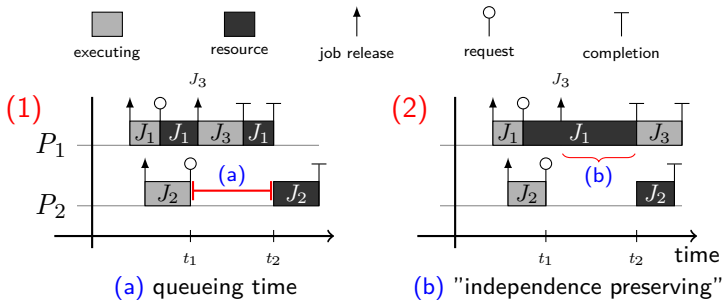
MrsP

Proposed
solution

Implementation

Experiments

Conclusion




- 2 processors and 3 tasks, $prio(J_i) > prio(J_y) \iff i > y$
- J_1 and J_2 share a global resource
- (1) J_1 remains preemptable
- (2) J_1 inhibits preemption



Suspension-based or spin-based protocol

A job, attempting to access a busy resource, will self-suspend (J_2) or will perform busy-wait (J_3)





Burns and Wellings [4] design a multiprocessor extension of PCP/SRP [1] with the aim of adapt a schedulability analysis to the protocol

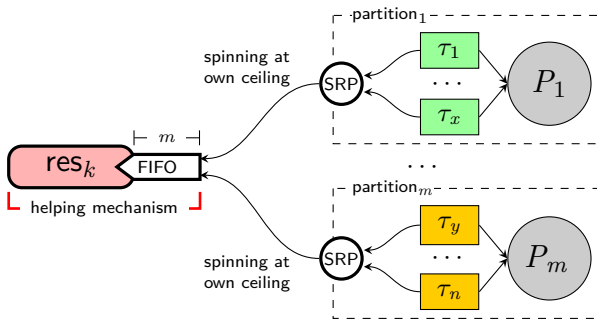
Response Time Analysis incorporating PCP/SRP

The parameter e_j reflects the contention for the resource (r):

$$e_j = |\text{map}(G(r))| \times c_j$$

$$R_i = C_i + \max\{e_j, \hat{b}\} + \sum_{\tau_j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$$

$$C_i = WCET_i + \sum_{r^j \in F(\tau_i)} n_i e_j$$



Protocol's properties

- It inherits the properties of PCP/SRP
- At most one job per processor requires the resource
- The length of the requests queue is at most $|map(G(r_j))|$
- At most e_j to gain the resource and to execute the critical section




MrsP

Runtime example

Sebastiano
Catellani

Introduction
MrsP
Proposed solution
Implementation
Experiments
Conclusion



- t_1 : J_1 's priority is raised: it gains access to r
- t_2 : J_3 's priority is raised: it starts spinning
- t_3 : J_2 is released and J_3 "helps" J_1



Proposed solution

Algorithm

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion

- 1) Each resource has a set of ceilings, one for each processor
- 2) An access request causes the rise of the job's priority and activates a local ceiling
- 3) The requests are queued and served in arrival order
- 4) A job executes, until resource's release, at the inherited priority
- 5) If preempted, the lock holder migrates to the first processor available

Key features

- Points 2 and 4 make MrsP independence-preserving
- Point 5 guarantees a limited waiting and blocking time



Implementation

Data structures

IMSP

Sebastiano
Catellani

Introduction


MrsP

Proposed
solution

Implementation

Experiments

Conclusion





Implementation

Queue management - 1

IMSP

Sebastiano
Catellani

Introduction

MrsP


Proposed
solution

Implementation

Experiments

Conclusion

Focused on managing the access requests queue



If preempted, the lock holder (J_x)

- 1 inherits the ceiling ($c_3 + 1$)
- 2 migrates to P_3
- 3 preempts J_y



Implementation

Queue management - 2

MrsP

Sebastiano
Catellani

Introduction

MrsP


Proposed
solution

Implementation


Experiments

Conclusion


The job will be re-queued in the *ready_queue*



The algorithm catches the operations that



(a) a processor becomes available




(b) add a new request to the queue



Implementation

Primitive: `mrsp_lock`



Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion



Implementation

Primitive: mrsp_unlock

MrsP

Sebastiano
Catellani

Introduction

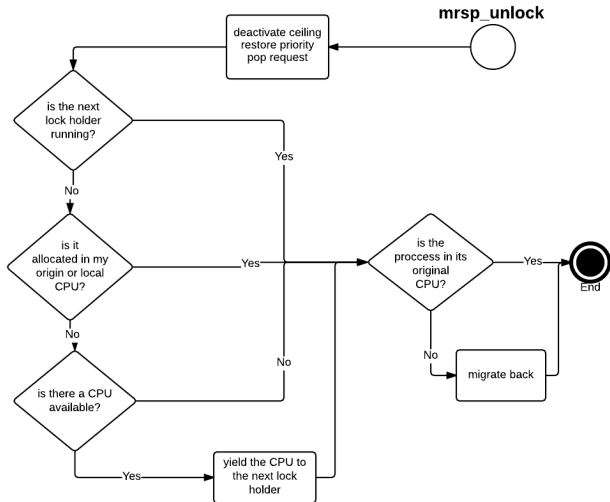
MrsP

Proposed
solution

Implementation

Experiments


Conclusion





Implementation

Primitive: pfp_schedule and finish_switch - 1



- t_1 : J_1 is marked for migration
- t_2 : J_4 's priority is lower than the local ceiling
- t_3 : default migration mechanism



Implementation

Primitive: pfp_schedule and finish_switch - 2

Sebastiano
Catellani

Introduction


MrsP

Proposed
solution

Implementation

Experiments

Conclusion



- t : J_2 completes and P_1 returns available



Experiments

Overview

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion

Experiment #1: Comparison among protocols

MrsP outperforms protocols based on simple ceiling or non preemption

Experiment #2: Sampling of the overheads

MrsP brings benefits at reasonable costs

Experiment #3: Absence of global resources

The protocol doesn't interfere with the scheduler



Experiments

Comparison among protocols - 1

MrsP

Sebastiano
Catellani

Introduction

MrsP

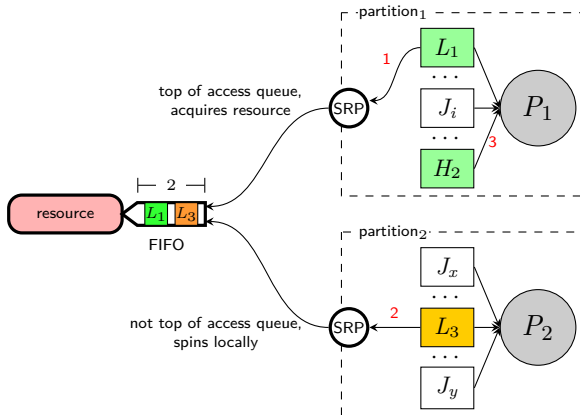
Proposed
solution

Implementation

Experiments

Conclusion

The experiment observes the **response times** of L_1 , H_2 and L_3 while varying the **critical section length** and the **WCET** of H_2





Experiments

Comparison among protocols - 2

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion




Figure: Response time of L_1



Experiments

Comparison among protocols - 3

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion




Figure: Response time of H_2



Experiments

Comparison among protocols - 4

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion




Figure: Response time of L_3



Experiments

Sampling of the overheads

MrsP

Sebastiano
Catellani

Introduction


MrsP

Proposed
solution


Implementation

Experiments


Conclusion



(a) `mrsp_lock`



(b) `mrsp_unlock`



(c) `finish_switch`



Experiments

MrsP without global resources

MrsP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion

The collected data show the same number of deadline miss




Figure: Number of *deadline miss*



Experiments

MrsP without global resources - pfp_schedule performance

MrsP

Sebastiano
Catellani

Introduction


MrsP

Proposed
solution


Implementation

Experiments


Conclusion



(a) pfp_schedule: Min



(b) pfp_schedule: Max



(c) pfp_schedule: Average



Conclusion

Future work - 1

The experiments underline how the system suffers in case of **migration**

A possible solution

- 1 Divide the resources into "long" and "short" (FMLP, [2])
- 2 Migration should only be caused when its benefits exceeds its cost

Supporting **nested resources**

- Allow a nested request from r_i to r_j only if $i < j$
- Groups lock (FMLP)
- A k-lock system (RNLP, [5])

Conclusion

Future work - 2

OMIP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution


Implementation

Experiments

Conclusion

Compare MrsP and the $O(m)$ Independence-preserving Protocol(OMIP, [3])

- Independence preserving and limited waiting time
- Allows migrations
- Design for cluster scheduler
- Suspension-based





References

IMRP

Sebastiano
Catellani

Introduction

MrsP

Proposed
solution

Implementation

Experiments

Conclusion



T. P. Baker.

Stack-based scheduling for realtime processes.
Real-Time Syst., 3(1):67–99, April 1991.



Aaron Block, Hennadiy Leontyev, Bjorn B. Brandenburg, and James H. Anderson.

A flexible real-time locking protocol for multiprocessors.
In Proceedings of the 13th IEEE International Conference on Embedded and Real-Time Computing Systems and Applications, RTCSA '07, pages 47–56, Washington, DC, USA, 2007. IEEE Computer Society.



B.B. Brandenburg.

A fully preemptive multiprocessor semaphore protocol for latency-sensitive real-time applications.
In Real-Time Systems (ECRTS), 2013 25th Euromicro Conference on, pages 292–302, July 2013.



A. Burns and A. J. Wellings.

A schedulability compatible multiprocessor resource sharing protocol – mrsp.
In Proceedings of the 2013 25th Euromicro Conference on Real-Time Systems, ECRTS '13, pages 282–291, Washington, DC, USA, 2013. IEEE Computer Society.



Bryan C. Ward and James H. Anderson.

Supporting nested locking in multiprocessor real-time systems.
In ECRTS, pages 223–232, 2012.